

Name: **SOLUTIONS** _____

netid: _____

CS 425/ECE 428 Distributed Systems
Mid-Term Exam 2

Total points: 50

Duration: 75 minutes

The exam includes 9 questions.

1. Consider a Chord peer-to-peer network consisting of 5 nodes that use 6-bit identifiers. The node identifiers are N12, N15, N29, N33 and N54.

(a) (3 points) Show the finger table at node N15 using the format below.


i	ft[i]
0	29
1	29
2	29
3	29
4	29
5	54

(b) (2 points) If node N15 wants to locate key K6, identify the node to which N15 will send its request message.

Node N54

2. (5 points) Assume the following: (i) node D below can sign its messages using its private key, (ii) the network contains at most 1 Byzantine faulty node, and (iii) all nodes know D's public key.

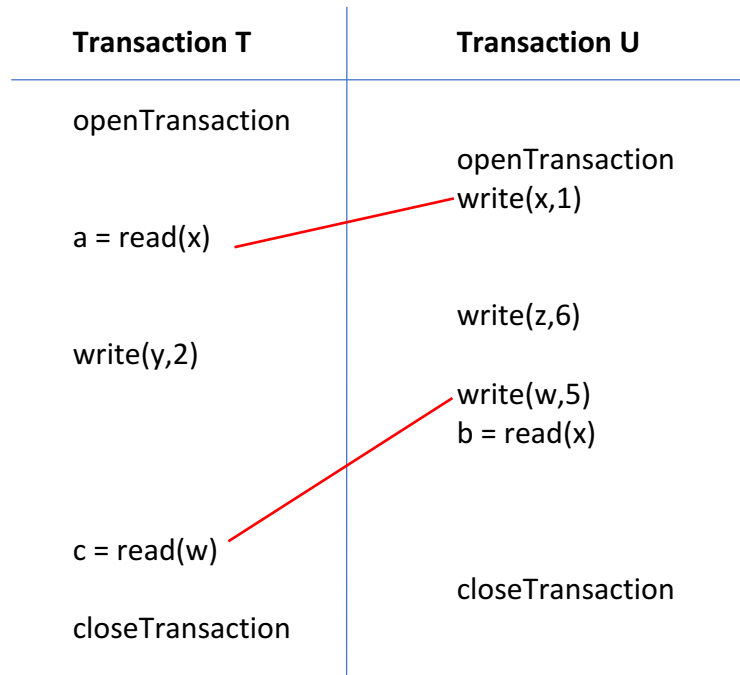
Answer True or False in each part below.



(a) Node B can always reliably receive messages sent by node D. TRUE

(b) Node F can always reliably receive messages sent by node D. FALSE

3. (a) (3 points) In the figure below, identify conflicting operations of transactions T and U.
- (b) (2 points) Is the interleaving of transactions T and U below *serially equivalent*?
 Answer YES or NO YES



4. (6 points) Does the shared memory algorithm presented below guarantee that a process that wants to enter the critical section will eventually enter the critical section? Assume that flag is initialize to 0. Answer Yes or No, and **justify your answer**.

Code for entry section:

1 Wait until flag = 0
 2 flag = 1

Code for exit section:

3 flag = flag - 1


NO.

Two processes may both find flag = 0 on line 1, and both proceed to line 2, and then both set flag = 1, and then both enter critical section.

When they both exit, the decrement flag by 1 in line 3, and the fina value of flag is now -1.

All processs wanting to enter the critical section subsequently will wait at line 1 indefinitely, since flag is now -1.

5. (5 points) The network below uses the link reversal algorithm to maintain routes to **destination B**. The figure shows the directions assigned to the links at a particular time.



Show how the link reversal algorithm will change directions of the links, if link BD now breaks. Present your answer in the form of a sequence of network states.

In the first step D reverses all its links.

After that, A reverses its link.

6. (6 points) Consider the Bakery algorithm for mutual exclusion (included in the attached handout).

Suppose that the process with the largest identifier is *Byzantine faulty*, and may write arbitrary values to the shared memory.

In this system, is it possible for the faulty process to prevent non-faulty processes from entering the critical section?

Justify your answer.

There are several possible answer.

One possibility is for a faulty process i to set $choosing[i]$ flag to true forever.

7. (6 points) This question relates to the Paxos algorithm discussed in class.

State True or False:

- (a) Paxos is guaranteed to perform correctly even if some of the acceptors are Byzantine faulty. FALSE
- (b) In an execution in which only crash failures occur, it is not possible for different learners to learn different chosen values. TRUE
- (c) If an acceptor responds to a *prepare* request with sequence number n , then no proposer will later send that acceptor a *prepare* request with sequence number less than n .
FALSE

8. (6 points) Determine whether the interleaving of operations of transactions T and U shown below can occur in each of the three cases below. In each case, answer YES or NO.

(a) Read-write locks are used with strict two-phase locking. YES

(b) Exclusive locks are used with two-phase locking. NO

(c) Exclusive locks are used with strict two-phase locking. NO

Transaction T	Transaction U
openTransaction a = read (x) write (w,2) d = read (y) closeTransaction	openTransaction write (y,5) c = read (x) closeTransaction

9. (6 points) The table below shows the interleaving of the operations performed by transactions T, U and V. Suppose that optimistic concurrency control with **backward** validation is used.

Transaction T	Transaction U	Transaction V
openTransaction write (z,9) a = read (v) write (y,7) write(q,3) closeTransaction	openTransaction write(v,8) f = read(w) write(w,8) write(p,2) closeTransaction	openTransaction c = read (v) write (w,5) closeTransaction

- (a) Will transaction U be aborted? Answer YES or NO, and **briefly explain why**.

YES.

Transaction V commits, and read set of U intersects with write set of V.

- (b) Will transaction T be aborted? Answer YES or NO, and **briefly explain why**.

NO.

Transaction V commits, and transaction U aborts, as noted above.

Read set of T does not intersect with write set of V.

Handout for Exam 2

Bakery Algorithm

Code for entry section:

```

Choosing[i] := true
Number[i] := max{Number[0], ..., Number[n-1]} + 1
Choosing[i] := false
for j := 0 to n-1 (except i) do
    wait until Choosing[j] = false
    wait until Number[j] = 0 or
        (Number[j,j] > (Number[i],i))
endfor

```

Code for exit section:

```

Number[i] := 0

```

2-Processor Mutex Algorithm

Code for entry section:

```

1 W[i] := 0
2 wait until W[1-i] = 0 or Priority = i
3 W[i] := 1
4 if (Priority = 1-i) then
5     if (W[1-i] = 1) then goto Line 1
6     else wait until (W[1-i] = 0)

```

Code for exit section:


```

7 Priority := 1-i
8 W[i] := 0

```

Deadlocks

- Necessary conditions for deadlocks
 - Non-shareable resources (locked objects)
 - No preemption on locks
 - Hold & Wait
 - Circular Wait (Wait-for graph)



Validation of Transactions

Backward validation of transaction T_v

```

boolean valid = true;
for (int  $T_i = startT_n + 1$ ;  $T_i \leq finishT_n$ ;  $T_i++$ ) {
    if (read set of  $T_v$  intersects write set of  $T_i$ ) valid = false;
}

```

Forward validation of transaction T_v


```

boolean valid = true;
for (int  $T_{id} = active1$ ;  $T_{id} \leq activeN$ ;  $T_{id}++$ ) {
    if (write set of  $T_v$  intersects read set of  $T_{id}$ ) valid = false;
}

```

Instructor's Guide for "Coulouris, Dollimore, Kindberg and Blair, Distributed Systems: Concepts and Design
Edn. 5
© Pearson Education 2012

Link Reversal Algorithm



Links are bi-directional
But algorithm imposes
logical directions on them

Maintain a directed acyclic
graph (DAG) for each
destination, with the destination
being the **only sink**

This DAG is for **destination
node D**

Peer pointers (2): finger tables

Finger Table at N80

